

NutriLabelAI - Complete Repository Guide

Project Name: NutriLabelAI

Repository: github.com/yvagula06/Senior-Design-2025

Description: AI-powered nutrition estimation system with semantic search, camera-based meal detection, mobile app, and comprehensive food database

Project Overview

NutriLabelAI is a production-ready nutrition estimation system that helps users get accurate nutrition information for any dish by combining semantic search with a comprehensive food database.

What It Does


Text-Based Label Generation: Enter a dish name like "chicken tikka masala" or "Big Mac" and get:


- ✓ Full FDA-style nutrition label: **calories, protein_g, carbs_g, fat_g** (required) plus up to 12 nullable micronutrients: **fiber_g, sugar_g, sodium_mg, potassium_mg, saturated_fat_g, trans_fat_g, cholesterol_mg, vitamin_a_mcg, vitamin_c_mg, vitamin_d_mcg, calcium_mg, iron_mg**
- ✓ Confidence score indicating reliability
- ✓ Best match from 2,400,000+ dishes in database
- ✓ Instant results via mobile app or API
- ✓ Optional meal log creation: when `device_id` is provided the response includes a `meal_log_id`

Camera-Based Meal Estimation: Take a photo of your meal and get:

- ✓ Automatic dish identification from image
- ✓ Volume/portion size estimation
- ✓ Calorie estimates with accuracy ranges
- ✓ Multiple estimation modes (depth, multi-angle, reference-based)

System Architecture





Core Features

- 🍽️ **2,400,000+ dishes** covering fast food, restaurants, home cooking, and international cuisines (USDA + Open Food Facts + curated seed data)
- 🔍 **Semantic search** - finds dishes even with typos, synonyms, or different phrasings
- 📷 **Camera-based estimation** - take a photo to get nutrition estimates with volume detection
- 📊 **Complete nutrition** - up to 16 FDA nutrients per dish (4 required macros + up to 12 nullable micronutrients: fiber, sugar, sodium, potassium, saturated fat, trans fat, cholesterol, vitamins A/C/D, calcium, iron)
- 🎯 **Confidence scoring** - tells you how reliable the match is (0.0-1.0)
- 📱 **Native mobile app** - iOS & Android via React Native + Expo
- ⚡ **Fast retrieval** - sub-100ms queries with pgvector HNSW indexing
- 🔄 **Portion scaling** - adjust nutrition based on target calories

🏗️ System Architecture

Three-Tier Architecture

1. Presentation Layer - Mobile App

- React Native + Expo for cross-platform support
- TypeScript for type safety
- React Native Paper for Material Design UI
- Expo Go for instant testing on physical devices


2. Application Layer - FastAPI Backend


- RESTful API with automatic OpenAPI documentation
- Service-oriented architecture (retrieval, scaling, confidence services)
- Sentence-transformers for embedding generation
- Pydantic for request/response validation

3. Data Layer - PostgreSQL + pgvector

- Relational database with vector extension — **7 canonical tables**: `users`, `dishes`, `dish_variants`, `meal_logs`, `vision_estimates`, `vision_feedback`, `user_portion_preferences`
- HNSW cosine index on `dish_variants.embedding` for sub-100ms semantic search
- Alembic for version-controlled schema migrations
- Foreign key relationships for data integrity (`SET NULL` on cross-table FKs to preserve history)

How It Works: Label Generation Flow





```

"nutrition": {
  "calories": 550,
  "protein_g": 25.5,
  "carbs_g": 45.2,
  "fat_g": 29.0,
  "fiber_g": 3.0,
  "sugar_g": 9.0,
  "sodium_mg": 1010,
  "potassium_mg": null,
  "saturated_fat_g": 10.0,
  "trans_fat_g": 1.0,
  "cholesterol_mg": 80,
  "vitamin_a_mcg": null,
  "vitamin_c_mg": null,
  "vitamin_d_mcg": null,
  "calcium_mg": null,
  "iron_mg": null
},
"confidence": 0.91,
"explanation": "Excellent match",
"meal_log_id": 101
}
NOTE: meal_log_id is non-null only
when device_id was in the request.

```

JSON response

Mobile App (Result Screen)

- Displays full nutrition label
- Shows confidence badge
- If device_id was sent, meal was auto-saved to meal_logs table.
- Client may also call POST /meal-logs explicitly to save.

Complete Directory Structure

Senior-Design-2025/

├──  mobile/	# React Native Mobile App
│ ├── src/	
│ │ ├── components/	# Reusable UI components
│ │ │ ├── Button.tsx	# Custom buttons
│ │ │ ├── Card.tsx	# Content cards
│ │ │ ├── TextInput.tsx	# Text inputs
│ │ │ ├── LoadingSpinner.tsx	# Loading indicator
│ │ │ ├── MacroPieChart.tsx	# Macro distribution chart
│ │ │ └── index.ts	# Component exports

```

├── Explore/                                # Explore-specific components
│   ├── CategoryHeader.tsx
│   ├── DishCard.tsx
│   └── index.ts
├── History/                                # History-specific components
│   ├── HistoryItemCard.tsx
│   └── index.ts
├── Label/                                  # Label-specific components
│   ├── ConfidenceBar.tsx
│   ├── DishSearchInput.tsx
│   ├── NutritionLabelCard.tsx
│   ├── VariantBottomSheet.tsx
│   ├── VariantDrawerButton.tsx
│   └── index.ts
├── Toast.tsx                              # Toast notification component
├── Profile/                                # Profile-specific components
│   ├── InfoCard.tsx
│   ├── SectionHeader.tsx
│   ├── SettingsItem.tsx
│   └── index.ts
├── Vision/                                # Vision-specific components
│   ├── AngleIndicator.tsx
│   ├── ARScanningOverlay.tsx
│   ├── CalorieRangeDisplay.tsx
│   ├── CameraGuide.tsx
│   ├── CaptureModeSelector.tsx
│   ├── DishPredictionList.tsx
│   ├── PlateReferenceSelector.tsx
│   ├── ReferenceObjectSelector.tsx
│   ├── SegmentationPreview.tsx
│   └── index.ts
├── config/                                # App configuration
│   └── fonts.ts                            # Font definitions
├── context/                                # State management
│   ├── FoodContext.tsx                    # Global food state
│   └── ThemeContext.tsx                   # Theme/dark-mode state
├── hooks/                                  # Custom React hooks
│   ├── index.ts                           # Hook exports
│   └── useDepthCamera.ts                   # Depth camera hook
├── native/                                # Native module bridges
│   └── DepthExtractor.ts                   # Native depth data extraction
├── navigation/                             # Navigation setup
│   ├── RootTabNavigator.tsx               # Root tab navigator
│   ├── BottomTabNavigator.tsx            # Bottom tab bar
│   ├── LabelStackNavigator.tsx           # Label flow stack
│   ├── ExploreStackNavigator.tsx         # Explore flow stack
│   ├── HistoryStackNavigator.tsx         # History flow stack
│   └── types.ts                            # Navigation type definitions

```

```

├── screens/                                # Main app screens
│   ├── Label/                             # Label generation screens
│   │   ├── LabelHomeScreen.tsx           # Input form
│   │   ├── LabelResultScreen.tsx         # Results display
│   │   └── BarcodeScannerScreen.tsx      # Barcode scan interface
│   ├── Vision/                           # Camera-based screens
│   │   ├── CameraCaptureScreen.tsx       # Camera interface
│   │   └── EstimationResultScreen.tsx     # Vision results
│   ├── History/                          # History screens
│   │   ├── HistoryListScreen.tsx         # List view
│   │   └── HistoryDetailScreen.tsx       # Detail view
│   ├── AddEntryScreen.tsx                # Manual food entry
│   ├── DailyConsumerScreen.tsx           # Daily tracking/history
│   ├── ExploreScreen.tsx                 # Browse dishes
│   ├── ProfileScreen.tsx                 # User profile
│   └── index.ts                           # Screen exports
├── services/                             # API communication
│   ├── api.ts                            # Axios client config
│   ├── barcode.ts                        # Barcode scanning service
│   ├── label.ts                          # Label API calls
│   ├── labelApi.ts                      # Label service
│   ├── visionApi.ts                     # Vision API calls
│   ├── DepthExtractor.ts                 # Depth data extraction service
│   └── storage.ts                        # AsyncStorage wrapper
├── theme/                                # Design system
│   ├── animations.ts                    # Animation definitions
│   ├── colors.ts                        # Color palette
│   ├── constants.ts                     # Theme constants
│   ├── typography.ts                    # Typography styles
│   └── index.ts                          # Theme exports
├── types/                                # TypeScript definitions
│   ├── label.ts                         # Label types
│   ├── nutrition.ts                     # Nutrition data types
│   └── vision.ts                         # Vision/camera types
├── assets/                               # Static assets
│   ├── fonts/                           # Custom fonts
│   └── images/                           # Images & icons
├── App.tsx                              # Root component
├── app.json                              # Expo config
├── package.json                          # Dependencies
├── tsconfig.json                         # TypeScript config
├── README.md                             # Mobile documentation
├── 🗝️ app/                                # FastAPI Backend
│   ├── api/                             # API route handlers
│   │   ├── __init__.py
│   │   └── dishes_router.py             # GET /dishes/search?q=&k= - Semantic dish
│   └── search
│       ├── label_router.py              # POST /label - Generate labels ★

```

	vision_router.py	# POST /vision/estimate - Camera estimation # POST /vision/feedback - Submit correction # GET /vision/personalization/{id} - User
profile		
		# GET /vision/feedback/stats - Aggregated
stats		
	feedback_router.py	# POST /meal-logs - Save meal log ★ # GET /meal-logs/{device_id} - Retrieve
history		
	core/	# Configuration
	__init__.py	
	settings.py	# Environment variables
	db/	# Database layer
	__init__.py	
	models.py	# SQLAlchemy ORM models ★ (7 tables) # - User, Dish, DishVariant # - MealLog, VisionEstimate # - VisionFeedback, UserPortionPreference
	session.py	# Database connection
	schemas/	# Pydantic schemas
	__init__.py	
	label.py	# Request/response models ★ # - LabelRequest # - LabelResponse
	vision.py	# Vision request/response models # - VisionRequest # - VisionResponse
	services/	# Business logic ★
	__init__.py	
	retrieval_service.py	# pgvector similarity search
	mixture_service.py	# Weighted aggregation
	scaling_service.py	# Calorie-based scaling
	rebalance_service.py	# Nutrient rebalancing
	confidence_service.py	# Confidence scoring
	vision_orchestrator.py	# Vision pipeline orchestration
	vision_api_client.py	# External vision API client
	vision_feedback_service.py	# Vision feedback tracking
	segmentation_service.py	# Dish segmentation from images
	volume_estimator.py	# Volume/portion size estimation
	depth_volume_estimator.py	# Depth-based volume estimation
	dish_classifier.py	# Dish type classification
	nutrition_mapper.py	# Volume to nutrition mapping
	vision_confidence_adapter.py	# Vision confidence scoring
	utils/	# Utilities
	__init__.py	
	embeddings.py	# Sentence-transformer loader
	static/	# Served static files (app-logo.png, etc.)


```

├── tests/                                # Unit tests
│   ├── test_confidence.py                # Confidence service tests
│   ├── test_health.py                    # Health endpoint tests
│   ├── test_label_flow.py                # End-to-end tests
│   ├── test_rebalance.py                 # Rebalance tests
│   ├── test_confidence_bounds.py          # Boundary tests
│   ├── test_end_to_end.py                # Integration tests
│   ├── test_label_router.py              # Router tests
│   ├── test_scaling_edge_cases.py         # Edge case tests
│   ├── test_vision_estimate_contract.py   # Vision API contract tests
│   └── test_vision_orchestrator_mocked.py # Vision orchestrator unit tests
├── __init__.py
├── main.py                               # FastAPI app entry point ★
├── 🗄️ alembic/                             # Database Migrations
│   ├── versions/
│   │   ├── 0001_init_schema_with_pgvector.py ★
│   │   │   # Creates ALL 7 tables: users, dishes, dish_variants,
│   │   │   # meal_logs, vision_estimates, vision_feedback,
│   │   │   # user_portion_preferences
│   │   │   # Installs pgvector extension + HNSW cosine index
│   │   │   # Drops any legacy tables before creating canonical schema
│   │   │   # Installs _update_updated_at trigger (users/dishes/variants)
│   │   └── 0002_add_vision_feedback_schema.py
│   │       # No-op placeholder (revision: 0002_noop)
│   │       # All tables live in 0001_canonical; keeps revision chain intact
│   └── env.py                            # Migration environment
├── 📂 data/                                # Datasets ★ (gitignored – see Database
section)
│   ├── seed_dishes.csv                   # Curated dishes (1 row)
│   ├── fastfood.csv                      # Fast food items (515 rows)
│   ├── usda_branded_foods.csv            # USDA database (full, ~2.4M entries)
│   ├── usda_branded_foods_reduced.csv    # USDA subset
│   ├── foundation/                       # USDA Foundation Foods CSV set
│   ├── sr_legacy/                         # USDA SR Legacy CSV set
│   └── survey/                           # USDA FNDDS survey foods CSV set
├── 📄 Implementation_Plans/               # Feature design documents
│   └── Camera_Functionality_Plan.md       # Camera pipeline architecture plan
├── 🎨 Poster/                             # Academic poster assets
│   ├── diagrams.py                       # Diagram generation script
│   ├── diagrams_copy.py                  # Diagram variant script
│   ├── architecture_diagram.png          # Architecture diagram (PNG)
│   ├── architecture_diagram.svg          # Architecture diagram (SVG)
│   ├── custom_pipeline.svg               # Custom pipeline graphic (SVG)
│   ├── pipeline_flow_diagram.png         # Pipeline flow (PNG)
│   ├── pipeline_flow_diagram.svg         # Pipeline flow (SVG)
│   ├── pipeline_vertical_zigzag.png      # Vertical pipeline layout (PNG)
│   └── pipeline_vertical_zigzag.svg       # Vertical pipeline layout (SVG)
├── 🤖 ml_models/                          # Trained Models

```

└─ neural_network_model.pth	# PyTorch MLP (archived)
└─  scripts/	# Data & Admin Scripts ★
└─ ingest_seed.py	# Load seed dishes
└─ embed_dishes.py	# Generate embeddings
└─ import_usda_dishes.py	# Populate dishes from USDA data ★★
└─ import_usda_fixed.py	# Fixed USDA import variant
└─ import_usda_whole_foods.py	# Import USDA whole foods subset
└─ import_fastfood.py	# Import fast food dataset
└─ import_openfoodfacts.py	# Import Open Food Facts data
└─ import_continuous.ps1	# Continuous import PowerShell script
└─ populate_db_simple.py	# Simplified database population
└─ reduce_dataset.py	# Preprocess CSVs
└─ inspect_db.py	# Database inspector
└─ query_db.py	# Test queries
└─ check_dishes.py	# Validate data
└─ insert_test_data.py	# Test records
└─ inspect_audit_logs.py	# Audit logs
└─ README_PREPROCESSING.md	# Preprocessing guide
└─  Jupyter Notebooks	# ML Experimentation
└─ DSA330_Nutrition_TextRegression.ipynb	# Text regression experiments
└─  Visualization Assets	
└─ pipeline_horizontal.png	# System pipeline diagram (PNG)
└─ pipeline_horizontal.svg	# System pipeline diagram (SVG)
└─ system_architecture.png	# Architecture overview (PNG)
└─ system_architecture.svg	# Architecture overview (SVG)
└─  Configuration Files	
└─ docker-compose.yml	# Docker orchestration ★
	# - nutrition_db (PostgreSQL)
	# - nutrition_api (FastAPI)
└─ Dockerfile	# Backend container image
└─ .dockerignore	# Docker ignore patterns
└─ pyproject.toml	# Python dependencies
└─ package.json	# Node.js dependencies
└─ tsconfig.json	# TypeScript config
└─ alembic.ini	# Alembic config
└─ Makefile	# Build automation
└─ .gitignore	# Git ignore patterns
└─  Documentation ★	
└─ README.md	# Main project overview
└─ QUICKSTART.md	# Complete setup walkthrough ★★
└─ QUICK_START_PHASE2.md	# Phase 2 quick start
└─ REPO_BREAKDOWN.md	# This file ★★
└─ REPO_BREAKDOWN.pdf	# PDF export of this file
└─ DATASET_PLAN.md	# Data acquisition strategy
└─ MOBILE_INTEGRATION.md	# Mobile-backend integration
└─ LABEL_ROUTER_API.md	# API documentation
└─ TROUBLESHOOTING.md	# Common issues & fixes
└─ CAMERA_SETUP_GUIDE.md	# Camera feature setup
└─ CAMERA_STATUS_AND_TODO.md	# Camera implementation status

└─	ANDROID_DEPTH_INTEGRATION.md	# Android depth sensor integration
└─	IOS_SETUP_PHASE2.md	# iOS phase 2 setup
└─	PHASE2_SETUP_GUIDE.md	# Phase 2 setup instructions
└─	PHASE2_IMPLEMENTATION_STEPS.md	# Phase 2 implementation plan
└─	PHASE2_STATUS.md	# Phase 2 completion status
└─	PHASE3_SETUP_GUIDE.md	# Phase 3 setup instructions
└─	schema.sql	# Database schema reference
└─	 Reports/	# Academic Reports & Deliverables
└─	NutriLabelAI_Final_Report.docx	# Final project report ★★
└─	NutriLabelAI_Final_Report_REVISED.md	# Revised final report (Markdown)
└─	NutriLabelAI_Final_Report_BACKUP.docx	# Final report backup
└─	NutriLabelAI_Final_Report_BACKUP2.docx	# Final report backup 2
└─	MIDTERM_REPORT.md	# Academic midterm report
└─	MIDTERM_REPORT.pdf	# PDF export of midterm report
└─	midterm_report.py	# Midterm report generation script
└─	final_report_extracted.txt	# Extracted text from final report
└─	fairposter.ipynb	# Poster/presentation notebook
└─	Senior Design Fall 2025 Proposal.pdf	# Project proposal
└─	 report_out/	# Generated Report Figures
└─	fig1_architecture.png	# Architecture figure
└─	fig2_dataflow.png	# Data flow figure
└─	fig3_vision_pipeline.png	# Vision pipeline figure
└─	midterm_report.docx	# Generated midterm report (Word)
└─	 Root-Level Utilities & Test Scripts	
└─	check_db_status.py	# Check database connection & row counts
└─	check_models.py	# Verify ML model files are present
└─	check-import-progress.ps1	# Monitor dish import progress
└─	debug_openai.py	# Debug OpenAI Vision API integration
└─	insert_sample_dishes.py	# Insert sample records for testing
└─	nutrilabel.py	# Standalone nutrition label utility
└─	seed_db.py	# Seed database with initial data
└─	setup-backend.ps1	# PowerShell backend setup script
└─	verify_nutrition_calc.py	# Verify nutrition calculation accuracy
└─	test_api_flow.py	# End-to-end API flow test
└─	test_nutrition_labels.py	# Nutrition label output tests
└─	test_openai_vision.py	# OpenAI Vision integration tests
└─	test_real_image.py	# Real image end-to-end test
└─	test_search.py	# Search functionality tests
└─	test_queries.ps1	# PowerShell query test script
└─	 Build Artifacts (gitignored)	
└─	__pycache__/	# Python bytecode
└─	node_modules/	# Node dependencies
└─	.expo/	# Expo build cache

★ = Essential for understanding the system

★★ = Start here for setup and overview

Key Components Explained

Backend Services (app/services/)

1. Retrieval Service ([retrieval_service.py](#))

```
# Finds similar dishes using vector similarity
query_embedding = model.encode("chicken tikka masala") # → 384-dim vector
results = db.execute(
    "SELECT * FROM dish_variants "
    "ORDER BY embedding <=> CAST(:query AS vector) LIMIT 5"
)
# Returns: [(Chicken Tikka Masala, 0.92), (Butter Chicken, 0.78), ...]
```

2. Mixture Service ([mixture_service.py](#))

```
# Combines multiple similar dishes using weighted average
weights = softmax([0.92, 0.78, 0.65]) # → [0.45, 0.32, 0.23]
aggregated_nutrition = sum(weight * dish.nutrition for weight, dish in
zip(weights, dishes))
# Prevents single dish from dominating (max 70% weight)
```

3. Scaling Service ([scaling_service.py](#))

```
# Adjusts nutrition based on target calories
scaling_factor = target_calories / base_calories # 600 / 540 = 1.11
scaled_protein = base_protein * scaling_factor # 23g * 1.11 = 25.5g
# Clamps factor between 0.1x and 10x for safety
```

4. Confidence Service ([confidence_service.py](#))

```
# Scores prediction reliability (0.0 to 1.0)
similarity_score = 0.92 # How close the match is
scaling_penalty = 0.05 # Penalty for large portion adjustments
consistency_bonus = 0.02 # Bonus for consistent top results
final_confidence = 0.92 - 0.05 + 0.02 = 0.89
```

Vision Services (app/services/)

1. Vision Orchestrator ([vision_orchestrator.py](#))

```
# Coordinates the entire vision pipeline
# Takes image(s) + metadata → Returns nutrition estimate
def estimate_from_images(images, depth_data=None, mode='reference_based'):
```

```
segments = segmentation_service.detect_dishes(images)
volume = volume_estimator.estimate_volume(segments, depth_data, mode)
dish_type = dish_classifier.classify(segments)
nutrition = nutrition_mapper.map_to_nutrition(dish_type, volume)
confidence = vision_confidence_adapter.score(segments, volume, dish_type)
return NutritionEstimate(nutrition, confidence)
```

2. Segmentation Service ([segmentation_service.py](#))

```
# Detects and isolates dishes from images
# Uses image processing to identify dish boundaries
# Returns segmented regions for volume estimation
```

3. Volume Estimator ([volume_estimator.py](#))

```
# Estimates portion size using three modes:
# - depth: Uses depth map + camera intrinsics (most accurate)
# - multi_angle: Triangulates from multiple photos
# - reference_based: Compares to standard portions in database
```

4. Dish Classifier ([dish_classifier.py](#))

```
# Identifies dish type from image features
# Returns dish predictions with confidence scores
# Integrated with semantic search for improved matching
```

5. Nutrition Mapper ([nutrition_mapper.py](#))

```
# Maps estimated volume to nutrition values
# Uses portion size + dish type → calorie/macro predictions
# Queries database for similar dishes and scales appropriately
```

6. Vision Confidence Adapter ([vision_confidence_adapter.py](#))

```
# Calculates reliability of vision-based estimates
# Factors: segmentation quality, dish recognition confidence, volume accuracy
# Returns accuracy score and uncertainty range
```

7. Depth Volume Estimator ([depth_volume_estimator.py](#))

```
# Specialized volume estimator using depth map data
# Uses Android/iOS depth sensor output + camera intrinsics
# Most accurate volume mode when depth hardware is available
```

8. Vision API Client ([vision_api_client.py](#))

```
# Client for communicating with external vision/AI services
# Handles image submission, response parsing, retry logic
```

9. Vision Feedback Service ([vision_feedback_service.py](#))

```
# Stores and processes user corrections to vision estimates
# Enables model improvement via logged feedback entries
```

Database Schema

7 tables total — all created in migration `0001_canonical`. Primary keys are `BIGSERIAL`. Cross-table foreign keys use `ON DELETE SET NULL` to preserve historical rows when a dish or user is removed. The one exception is `dish_variants` → `dishes` which uses `ON DELETE CASCADE`.

users table — Identity anchor; one row per device (no login required)

```
CREATE TABLE users (
  id          BIGSERIAL PRIMARY KEY,
  device_id   VARCHAR(255) NOT NULL UNIQUE,    -- mobile UUID on first install
  created_at  TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
  updated_at  TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP -- via trigger
);
CREATE INDEX idx_users_device_id ON users (device_id);
```

dishes table — Canonical nutrition knowledge base (2,400,000+ rows)

All nutrient values are per 100 g unless `serving_size_g` overrides the reference. The 4 macro columns are always required; the 12 micronutrient columns are nullable because source data may be incomplete.

```
CREATE TABLE dishes (
  id          BIGSERIAL PRIMARY KEY,
  name        VARCHAR(255) NOT NULL,
  -- Core macros (non-nullable)
  calories    NUMERIC(8,2) NOT NULL CHECK (calories >= 0),
  protein_g   NUMERIC(8,2) NOT NULL CHECK (protein_g >= 0),
  fat_g       NUMERIC(8,2) NOT NULL CHECK (fat_g >= 0),
```

```

carbs_g          NUMERIC(8,2) NOT NULL CHECK (carbs_g >= 0),
-- FDA micronutrients (nullable)
fiber_g          NUMERIC(8,2),
sugar_g          NUMERIC(8,2),
sodium_mg        NUMERIC(8,2),
potassium_mg     NUMERIC(8,2),
saturated_fat_g  NUMERIC(8,2),
trans_fat_g      NUMERIC(8,2),
cholesterol_mg   NUMERIC(8,2),
vitamin_a_mcg    NUMERIC(8,2),
vitamin_c_mg     NUMERIC(8,2),
vitamin_d_mcg    NUMERIC(8,2),
calcium_mg       NUMERIC(8,2),
iron_mg          NUMERIC(8,2),
-- Serving reference
serving_size_g   NUMERIC(8,2),
serving_size_unit VARCHAR(50) DEFAULT 'g',
-- Metadata
category_name    VARCHAR(100),
data_source      VARCHAR(100),
confidence_score NUMERIC(4,3),
is_active        BOOLEAN NOT NULL DEFAULT TRUE,
version          INTEGER NOT NULL DEFAULT 1,
created_at       TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
updated_at       TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP
);
CREATE INDEX idx_dishes_name          ON dishes (name);
CREATE INDEX idx_dishes_data_source  ON dishes (data_source);
CREATE INDEX idx_dishes_active       ON dishes (is_active) WHERE is_active = TRUE;

```

dish_variants table — Searchable text aliases with 384-dim embeddings (2,400,000+ rows)

```

CREATE TABLE dish_variants (
  id          BIGSERIAL PRIMARY KEY,
  dish_id     BIGINT NOT NULL REFERENCES dishes(id) ON DELETE CASCADE,
  variant_text TEXT NOT NULL,
  embedding   VECTOR(384) NOT NULL,    -- sentence-transformers/all-MiniLM-
L6-v2
  variant_type VARCHAR(50),
  language_code VARCHAR(10) NOT NULL DEFAULT 'en',
  search_count INTEGER NOT NULL DEFAULT 0,
  last_searched_at TIMESTAMPTZ,
  created_at       TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
  updated_at       TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
  UNIQUE (dish_id, variant_text)
);
CREATE INDEX idx_dish_variants_dish_id      ON dish_variants (dish_id);
CREATE INDEX idx_dish_variants_variant_text ON dish_variants (variant_text);
-- HNSW cosine index – used by retrieval_service.py (<=> operator)
CREATE INDEX idx_dish_variants_embedding_hnsw

```

```
ON dish_variants USING hnsw (embedding vector_cosine_ops)
WITH (m = 16, ef_construction = 64);
```

Relationship: 1 dish → many variants

- "McDonald's Big Mac" → variants: ["big mac", "mcdonald's big mac", "mcdonalds big mac"]
- Each variant has its own 384-dimensional embedding for semantic retrieval

meal_logs table — Single source of truth for confirmed user meal entries

Stores a frozen **nutrition_label** JSONB snapshot so historical records are never altered by future model updates. Created automatically by **POST /label** (when **device_id** is provided) or explicitly via **POST /meal-logs**.

```
CREATE TABLE meal_logs (
  id BIGSERIAL PRIMARY KEY,
  user_id BIGINT REFERENCES users(id) ON DELETE SET NULL,
  dish_id BIGINT REFERENCES dishes(id) ON DELETE SET NULL,
  vision_estimate_id BIGINT REFERENCES vision_estimates(id) ON DELETE SET NULL,
  entry_source VARCHAR(20) NOT NULL CHECK (entry_source IN
('manual', 'camera')),
  logged_dish_name VARCHAR(255) NOT NULL, -- denormalised; survives dish
deletion
  logged_calories NUMERIC(8,2) NOT NULL CHECK (logged_calories >= 0),
  nutrition_label JSONB, -- full FDA label snapshot
(nullable)
  serving_multiplier NUMERIC(6,3) NOT NULL DEFAULT 1.0,
  match_confidence NUMERIC(4,3),
  model_version VARCHAR(50),
  logged_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
  created_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP
);
CREATE INDEX idx_meal_logs_user_id ON meal_logs (user_id);
CREATE INDEX idx_meal_logs_logged_at ON meal_logs (logged_at);
CREATE INDEX idx_meal_logs_dish_id ON meal_logs (dish_id);
```

API endpoints (feedback_router.py):

- **POST /meal-logs** — saves a new entry (manual or camera-confirmed)
- **GET /meal-logs/{device_id}** — returns up to 50 most recent entries for a device

vision_estimates table — Audit log for every camera-pipeline prediction

Images are never stored as base64; only storage path/key references are kept.


```

CREATE TABLE vision_estimates (
  id BIGSERIAL PRIMARY KEY,
  user_id BIGINT REFERENCES users(id) ON DELETE SET NULL,
  capture_mode VARCHAR(20) NOT NULL
  CHECK (capture_mode IN
('depth','multi_angle','single')),
  device_type VARCHAR(10),
  num_images INTEGER NOT NULL,
  has_depth_data BOOLEAN NOT NULL DEFAULT FALSE,
  image_storage_keys JSONB,          -- paths/keys only; no base64
  depth_storage_key VARCHAR(500),
  predicted_dish_id BIGINT REFERENCES dishes(id) ON DELETE SET NULL,
  predicted_dish_name VARCHAR(255) NOT NULL,
  predicted_confidence NUMERIC(4,3),
  alternative_dishes JSONB,          -- [{dish_id, dish_name, confidence}],
  ...]
  estimation_mode VARCHAR(20) NOT NULL,
  volume_ml NUMERIC(8,2),
  volume_confidence NUMERIC(4,3),
  calorie_estimate NUMERIC(8,2) NOT NULL CHECK (calorie_estimate >= 0),
  calorie_range_min NUMERIC(8,2),
  calorie_range_max NUMERIC(8,2),
  classifier_version VARCHAR(50),
  segmentation_version VARCHAR(50),
  volume_estimator_version VARCHAR(50),
  processing_time_ms INTEGER,
  created_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP
);
CREATE INDEX idx_vision_estimates_user_id ON vision_estimates (user_id);
CREATE INDEX idx_vision_estimates_created_at ON vision_estimates (created_at);
CREATE INDEX idx_vision_estimates_dish_id ON vision_estimates
(predicted_dish_id);

```

vision_feedback table — User corrections on vision estimates (drives personalization)

```

CREATE TABLE vision_feedback (
  id BIGSERIAL PRIMARY KEY,
  vision_estimate_id BIGINT REFERENCES vision_estimates(id) ON DELETE SET NULL,
  user_id BIGINT REFERENCES users(id) ON DELETE SET NULL,
  corrected_dish_id BIGINT REFERENCES dishes(id) ON DELETE SET NULL,
  feedback_type VARCHAR(30) NOT NULL
  CHECK (feedback_type IN
('confirmed','corrected_dish','corrected_portion','quick_correction')),
  confirmed_dish_name VARCHAR(255),
  portion_adjustment NUMERIC(6,3),
  plate_size VARCHAR(30),
  quick_feedback VARCHAR(20),
  corrected_calories NUMERIC(8,2),

```

```

    notes                TEXT,
    created_at            TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP
);
CREATE INDEX idx_vision_feedback_estimate_id ON vision_feedback
(vision_estimate_id);
CREATE INDEX idx_vision_feedback_user_id    ON vision_feedback (user_id);

```

API endpoints (vision_router.py):

- **POST** /vision/feedback — submit correction on a vision estimate
- **GET** /vision/feedback/stats — aggregated feedback statistics
- **GET** /vision/personalization/{user_id} — retrieve personalization profile

user_portion_preferences table — Per-user personalization profile (one row per user)

Derived from **vision_feedback** history. Updated by **VisionFeedbackService** on every feedback write.

```

CREATE TABLE user_portion_preferences (
    id                BIGSERIAL PRIMARY KEY,
    user_id           BIGINT NOT NULL UNIQUE REFERENCES users(id) ON DELETE
CASCADE,
    avg_portion_factor NUMERIC(6,3) NOT NULL DEFAULT 1.0,
    feedback_count     INTEGER NOT NULL DEFAULT 0,
    confidence_score   NUMERIC(4,3) NOT NULL DEFAULT 0.0,
    dish_preferences   JSONB,      -- {dish_id_str: {avg_factor, count}}
    category_preferences JSONB,    -- {category: {avg_factor, count}}
    last_updated       TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
    created_at         TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP
);
CREATE INDEX idx_user_portion_preferences_user_id ON user_portion_preferences
(user_id);

```

Relationship summary:

```

users (1) ————— (N) meal_logs
users (1) ————— (N) vision_estimates
users (1) ————— (N) vision_feedback
users (1) ————— (1) user_portion_preferences
dishes (1) ————— (N) dish_variants [CASCADE DELETE]
dishes (1) ————— (N) meal_logs      [SET NULL on delete]
dishes (1) ————— (N) vision_estimates [SET NULL]
dishes (1) ————— (N) vision_feedback [SET NULL]
vision_estimates (1) — (N) vision_feedback
vision_estimates (1) — (0..1) meal_logs

```

Navigation Structure:

```
App.tsx
├── RootTabNavigator
│   ├── Label Tab (LabelStackNavigator)
│   │   ├── LabelHomeScreen (text input form)
│   │   ├── LabelResultScreen (nutrition card)
│   │   └── BarcodeScannerScreen (barcode scan interface)
│   ├── Camera Tab
│   │   ├── CameraCaptureScreen (camera interface)
│   │   └── EstimationResultScreen (vision results)
│   ├── History Tab (HistoryStackNavigator)
│   │   ├── DailyConsumerScreen (daily food tracking)
│   │   ├── AddEntryScreen (manual food entry)
│   │   ├── HistoryListScreen (meal history list)
│   │   └── HistoryDetailScreen (detailed view)
│   ├── Explore Tab (ExploreStackNavigator)
│   └── Profile Tab (settings)
```

State Management (FoodContext + ThemeContext):

```
// FoodContext – global food entries, daily totals, and streaks
const FoodContext = {
  foodEntries: [], // All FoodEntry[] objects
  calorieGoal: 2000, // User's daily calorie goal
  macroGoals: {...}, // Protein/carbs/fat goals
  streakDays: 0, // Consecutive logging streak
  recentEntries: [], // Last 5 unique dish names
  addLabelEntry: (entry) => {}, // Add from label/vision result
  deleteFoodEntry: (id) => {}, // Delete by ID
  clearAllData: () => {}, // Reset all entries
  setCalorieGoal: (n) => {}, // Update calorie goal
  setMacroGoals: (goals) => {}, // Update macro goals
  getTodayEntries: () => [], // Filter to today's entries
  getTotals: () => ({}), // Sum today's macros
};

// ThemeContext – dark/light mode preference with accent colours
const ThemeContext = {
  isDark: false, // Current theme mode
  accentName: 'amber', // Active accent: amber|emerald|sky|purple|rose
  toggleTheme: () => {}, // Switch between dark/light
  setAccent: (name) => {}, // Change accent colour
  colors: { ... } // Active colour palette (buildColors(isDark,
  accent))
};
```

API Service Configuration:

```
// Detects device type and configures API URL
const getBaseUrl = () => {
  const isDevice = Device.isDevice; // expo-device library

  if (Platform.OS === 'ios' && isDevice) {
    return 'http://192.168.1.191:8000'; // Physical iOS: LAN IP
  }
  if (Platform.OS === 'android' && !isDevice) {
    return 'http://10.0.2.2:8000'; // Android emulator: special alias
  }
  return 'http://localhost:8000'; // Default: localhost
};
```

❖ History & Meal Log — Implementation Status

What meal_logs Stores

Every confirmed meal entry lands in the `meal_logs` table. The row captures a **frozen snapshot** of the full nutrition label at the time of logging (`nutrition_label JSONB`), so historical records are never silently changed by future model or data updates. Key columns:

Column	Type	Notes
entry_source	VARCHAR(20)	'manual' (text search path) or 'camera' (vision path)
logged_dish_name	VARCHAR(255)	Denormalised — preserved even if dish is later deleted
logged_calories	NUMERIC	Calories at the time of logging, after any portion scaling
nutrition_label	JSONB	Full FDA label snapshot — all 16 nutrient fields
serving_multiplier	NUMERIC	Scaling factor applied (1.0 = no scaling)
vision_estimate_id	BIGINT FK	Links to vision_estimates row when source is camera
dish_id	BIGINT FK	Links to dishes row (SET NULL if dish later deleted)
user_id	BIGINT FK	Links to users row via device_id lookup

Two Paths That Create Meal Logs

Path A — Manual text label (via POST /label)

When the mobile client sends `device_id` in the `LabelRequest`, the router automatically writes a `meal_logs` row with `entry_source = 'manual'` and returns the new `meal_log_id` in the response. No second API call is required. If `device_id` is omitted (preview mode), no row is written.

Path B — Explicit save (via POST /meal-logs)

The client can always call `POST /meal-logs` directly with full control over all fields. This is intended for:

- Confirming a camera estimate the user wants to record
- Manual entry from `AddEntryScreen` without invoking the label pipeline

Camera Estimates and `vision_estimates`

Every call to `POST /vision/estimate` writes a row to `vision_estimates` regardless of whether the user saves the meal. The row contains the raw prediction output (dish name, confidence, calorie range, volume, model versions). A `meal_logs` row is created only when the user confirms the camera result, linking `meal_logs.vision_estimate_id` back to the estimate.

What Is Fully Implemented

Feature	Status	Backend endpoint
Save manual label as meal log	✓ Implemented	Auto-write in <code>POST /label</code> when <code>device_id</code> provided
Explicit meal log save	✓ Implemented	<code>POST /meal-logs</code>
Retrieve recent history	✓ Implemented	<code>GET /meal-logs/{device_id}</code> (latest 50 entries)
Camera estimate recording	✓ Implemented	Auto-write in <code>POST /vision/estimate</code>
Vision feedback submission	✓ Implemented	<code>POST /vision/feedback</code>
Personalization profile	✓ Implemented	<code>GET /vision/personalization/{user_id}</code>

What Is Partial or In Progress

Feature	Status	Notes
History filtering by date range	🚧 In Progress	<code>GET /meal-logs</code> returns newest 50; no date filter yet
Pagination of meal history	🚧 In Progress	No cursor/page param yet
Daily calorie totals API	🚧 In Progress	Aggregation not yet in backend; FoodContext handles client-side
History detail screen (mobile)	✓ Implemented	Full nutrition label displayed with real entry data; delete works
Camera-confirmed meal log (mobile)	✓ Implemented	<code>EstimationResultScreen.tsx</code> confirm flow wired to <code>addLabelEntry</code> in FoodContext

🔍🚀 Getting Started

Prerequisites

- Docker Desktop (for backend)
- Node.js 18+ (for mobile app)
- Expo Go app on your phone

Complete Setup (5 minutes)

1. Clone repository

```
git clone https://github.com/yvagula06/Senior-Design-2025.git
cd Senior-Design-2025
```

2. Start backend

```
docker-compose up -d --build          # Start containers
docker-compose exec api alembic upgrade head # Create schema
```

3. Populate database

⚠ The full database (~2.4M dishes) is not stored in git (datasets are gitignored due to size). See the **Database Distribution** section below for how to get the pre-populated dump.

If starting fresh with a small dataset for testing:

```
docker exec -it nutrition_api python /app/scripts/import_fastfood.py      #
~515 fast food items
docker exec -it nutrition_api python /app/scripts/ingest_seed.py          # 1
seed dish
```

For the full USDA dataset (takes 20-60 min):

```
docker exec -it nutrition_api python /app/scripts/import_usda_dishes.py
docker exec -it nutrition_api python /app/scripts/import_openfoodfacts.py
```

Database Distribution (pg_dump)

The production database (2.4M dishes + embeddings) is distributed as a **pg_restore**-compatible dump. Share via Google Drive / OneDrive and restore with:

```
# 1. Start the db container
docker-compose up -d db

# 2. Copy dump file into container
docker cp nutrition_dump.dump nutrition_db:/tmp/
```

```
# 3. Restore
docker exec nutrition_db pg_restore -U postgres -d nutrition
/tmp/nutrition_dump.dump
```

4. Configure mobile app

```
# Find your IP address (Windows)
ipconfig # Look for IPv4 Address

# Update mobile/src/services/api.ts with your IP
# Line 11: return 'http://YOUR_IP_ADDRESS:8000';
```

5. Start mobile app

```
cd mobile
npm install
npm start
```

6. Open on phone

- Scan QR code with Expo Go
- Test text search: "pizza" or "Big Mac"
- Test camera: take a photo of a meal

See [QUICKSTART.md](#) for detailed instructions.



Database Contents

Current Database (2,406,474 dishes, 2,406,474 variants)

Data Sources:

- [seed_dishes.csv](#) — 1 curated seed dish
- [fastfood.csv](#) — ~515 restaurant items
- USDA Branded Foods — majority of entries (~2.4M)
- USDA Foundation Foods, SR Legacy, Survey (FNDDS)
- Open Food Facts (supplemental)

Coverage:

- ☒ McDonald's, Burger King, Subway, Taco Bell, Pizza Hut, KFC (full menus)
- ☒ USDA branded & packaged foods
- ☒ USDA whole foods & ingredients
- ☒ International cuisines
- ☒ Home cooking ingredients

⚠️ All CSV data files are gitignored. The populated database is shared as a `pg_dump` file. Contact the team or check the shared drive for `nutrition_dump.dump`.

Variant Examples:

```
Dish: "McDonald's Big Mac"
Variants:
- "big mac"
- "mcdonald's big mac"
- "mcdonalds big mac"

Dish: "KFC Original Recipe Chicken"
Variants:
- "kfc chicken"
- "original recipe chicken"
- "kfc original chicken"
```

Future Expansion (Planned)

See [DATASET_PLAN.md](#) for strategy to expand to 800+ dishes covering:

- Home-style cooking
- International cuisines
- Restaurant preparations
- Meal combinations

 Tech Stack Summary

Frontend (Mobile App)

Component	Technology	Purpose
Framework	React Native 0.78	Cross-platform mobile development
Runtime	Expo SDK 54	Development tools & native APIs
Language	TypeScript 5.9	Type-safe JavaScript
UI Library	React Native Paper	Material Design components
Navigation	React Navigation 7	Tab & drawer navigation
State	Context API	Global state management
HTTP Client	Axios	API requests
Device Detection	expo-device	Physical device detection
Storage	AsyncStorage	Local data persistence

Backend (API Server)

Component	Technology	Purpose
Framework	FastAPI 0.115	High-performance Python API
Server	Uvicorn	ASGI server
Language	Python 3.11+	Backend logic
ORM	SQLAlchemy 2.0	Database interactions
Migrations	Alembic 1.13	Schema version control
Validation	Pydantic 2.7	Request/response schemas
ML Library	sentence-transformers 3.0	Embedding generation
Vector Store	pgvector 0.2.5	PostgreSQL extension

Database

Component	Technology	Purpose
RDBMS	PostgreSQL 16	Relational database
Vector Extension	pgvector	Vector similarity search
Embedding Model	all-MiniLM-L6-v2	384-dimensional embeddings
Index Type	HNSW	Fast approximate nearest neighbor
Connection Pool	psycopg 3.2	Database connections

Machine Learning

Component	Technology	Purpose
Embedding	Sentence-transformers	Text → vector conversion
Deep Learning	PyTorch 2.x	Neural network framework
Classical ML	Scikit-learn	Linear regression, preprocessing
Numerical	NumPy, SciPy	Matrix operations
Notebooks	Jupyter	ML experimentation


DevOps & Tools

Component	Technology	Purpose
Containerization	Docker 24+	Application packaging
Orchestration	Docker Compose	Multi-container management
Version Control	Git + GitHub	Source code management
Testing	pytest	Unit & integration tests

Component	Technology	Purpose
Environment	Conda	Python environment management

Key Workflows

1. Label Generation (User → Mobile → Backend → Database)



▼

MIXTURE SERVICE

- Applies softmax to similarity scores
- Weights: [0.45, 0.28, 0.16, 0.08, 0.03]
- Calculates weighted average of nutrition:
calories = $0.45 \times 520 + 0.28 \times 580 + 0.16 \times 490 + \dots$
protein = $0.45 \times 18 + 0.28 \times 22 + 0.16 \times 16 + \dots$
- Prevents single dish from dominating (max 70% weight)

| Base nutrition: {calories: 540, protein: 19.2, ...}

▼

SCALING SERVICE

- Compares target (650) vs base (540) calories
- Calculates scaling factor: $650/540 = 1.20$
- Scales all nutrients proportionally:
protein: $19.2 \times 1.20 = 23.0\text{g}$
carbs: $45.0 \times 1.20 = 54.0\text{g}$
fat: $28.0 \times 1.20 = 33.6\text{g}$
- Clamps scaling factor (0.1x to 10x)

| Scaled nutrition

▼

CONFIDENCE SERVICE

- Base confidence from similarity: 0.89
- Scaling penalty: 1.20 scaling $\rightarrow$ -0.05 penalty
- Consistency bonus: top 3 all pasta $\rightarrow$ +0.02
- Final confidence: $0.89 - 0.05 + 0.02 = 0.86$
- Classification: "High" (≥ 0.7)

▼

API RESPONSE (LabelResponse)

```
{
  "matched_dish": "Fettuccine Alfredo",
  "nutrition": {
    "calories": 650.0,      "protein_g": 23.0,
    "carbs_g": 54.0,       "fat_g": 33.6,
    "fiber_g": 2.4,        "sugar_g": 3.6,
    "sodium_mg": 720,      "potassium_mg": 340,
    "saturated_fat_g": 12.0, "trans_fat_g": null,
    "cholesterol_mg": 55,   "vitamin_a_mcg": null,
    "vitamin_c_mg": null,   "vitamin_d_mcg": null,
    "calcium_mg": null,     "iron_mg": null
  },
  "confidence": 0.86,
  "explanation": "Strong match with consistent candidates",
  "meal_log_id": 42
}
```

| JSON response (200ms total)

MOBILE APP (LabelResultScreen.tsx)

- Hides loading spinner
- Displays nutrition card:

Fettuccine Alfredo [86% ✓]

Calories: 650 Protein: 23.0g

Carbs: 54.0g Fat: 33.6g

Fiber: 2.4g Sugar: 3.6g

Sodium: 720mg Potassium: 340mg

Sat.Fat: 12g Cholesterol: 55mg

(vitamins shown when data available)

- User can save to history
(meal_log_id in the response confirms auto-save;
or client calls POST /meal-logs explicitly)

2. Database Population Workflow

Clone Repository



Start Docker Containers

| docker-compose up -d --build



Create Database Schema

| docker-compose exec api alembic upgrade head
| Creates all 7 tables: users, dishes, dish_variants,
| meal_logs, vision_estimates, vision_feedback,
| user_portion_preferences
| Installs pgvector extension + HNSW cosine index



Populate Database (choose one path)

— OPTION A: Restore from pg_dump (recommended – fastest)

| docker cp nutrition_dump.dump nutrition_db:/tmp/
| docker exec nutrition_db pg_restore -U postgres \
| -d nutrition /tmp/nutrition_dump.dump
| → Loads 2.4M dishes + embeddings in minutes

— OPTION B: Run ingestion scripts from scratch

| docker exec -it nutrition_api python /app/scripts/import_fastfood.py
| docker exec -it nutrition_api python
/app/scripts/import_usda_dishes.py
| docker exec -it nutrition_api python
/app/scripts/import_openfoodfacts.py
| → Takes 20-60 min; generates all embeddings live



Database Ready ✓

- 2,406,474 dishes
- 2,406,474 variants with embeddings
- HNSW indexes built
- Ready for similarity search

3. Development Workflow

Developer Workflow:

1. BACKEND CHANGES

- └─ Edit Python code in app/
- └─ Docker auto-reloads (volume mount)
- └─ Test: `curl http://localhost:8000/health`
- └─ Check logs: `docker logs nutrition_api`

2. MOBILE CHANGES

- └─ Edit TypeScript in mobile/src/
- └─ Expo hot-reloads automatically
- └─ Shake phone → Developer menu
- └─ Press 'r' in terminal to reload

3. DATABASE CHANGES

- └─ Create migration: `alembic revision --autogenerate -m "message"`
- └─ Review: `alembic/versions/XXXX_message.py`
- └─ Apply: `docker-compose exec api alembic upgrade head`
- └─ Rollback: `docker-compose exec api alembic downgrade -1`

4. TESTING

- └─ Backend: `docker exec nutrition_api pytest app/tests/`
- └─ Mobile: `cd mobile && npm test`
- └─ Integration: Test end-to-end with Expo Go

5. DEPLOYMENT

- └─ Commit changes: `git add . && git commit -m "message"`
- └─ Push to GitHub: `git push origin main`
- └─ Deploy: (production setup TBD)



Testing

Test Suite Organization

Location: `app/tests/`

Test Files:

1. **test_health.py** - API health check endpoint
2. **test_label_flow.py** - End-to-end label generation
3. **test_confidence.py** - Confidence scoring logic

4. **test_rebalance.py** - Portion adjustment
5. **test_confidence_bounds.py** - Boundary conditions
6. **test_end_to_end.py** - Full pipeline integration
7. **test_label_router.py** - Router validation
8. **test_scaling_edge_cases.py** - Edge case handling
9. **test_vision_estimate_contract.py** - Vision API contract tests
10. **test_vision_orchestrator_mocked.py** - Vision orchestrator unit tests

Running Tests:

```
# All tests
docker exec nutrition_api pytest app/tests/ -v

# Specific test file
docker exec nutrition_api pytest app/tests/test_label_flow.py -v

# With coverage
docker exec nutrition_api pytest app/tests/ --cov=app --cov-report=html
```

Test Framework: pytest with async support

Coverage: Core business logic + API endpoints

Dependencies

Backend (Python) - pyproject.toml

```
[tool.poetry.dependencies]
python = "^3.11"
fastapi = ">=0.115.0"
uvicorn = {extras = ["standard"], version = ">=0.30.0"}
sqlalchemy = ">=2.0.0"
alembic = ">=1.13.0"
psycopg = {extras = ["binary", "pool"], version = ">=3.2.0"}
pgvector = ">=0.2.5"
sentence-transformers = ">=3.0.0"
torch = ">=2.0.0"
numpy = ">=1.26.0"
scipy = ">=1.12.0"
pydantic = ">=2.7.0"
python-multipart = ">=0.0.9"
```

Mobile (Node.js) - mobile/package.json

```
{
  "dependencies": {
    "expo": "^54.0.25",
```

```
"react": "18.3.1",
"react-native": "0.78.0",
"typescript": "~5.9.2",
"@react-navigation/native": "^7.0.14",
"@react-navigation/bottom-tabs": "^7.1.8",
"@react-navigation/drawer": "^7.0.8",
"react-native-paper": "^5.15.1",
"axios": "^1.7.9",
"expo-device": "^8.0.10",
"@expo/vector-icons": "^15.0.2"
}
}
```

Environment Setup

Environment Variables

Backend (.env or docker-compose.yml):

```
DATABASE_URL=postgresql+psycopg://postgres:postgres@db:5432/nutrition
EMBEDDING_MODEL=sentence-transformers/all-MiniLM-L6-v2
TOP_K=5
SIM_THRESHOLD=0.45
```

Mobile (api.ts):

```
// Configured in mobile/src/services/api.ts
const BASE_URL = 'http://YOUR_IP:8000'; // Update with your IP
```

Development Environment

Backend:

- Python 3.11+ (Conda environment recommended)
- Docker Desktop for containerization
- PostgreSQL 16 with pgvector (via Docker)

Mobile:

- Node.js 18+
- Expo CLI (installed via npm)
- Expo Go app on phone (for testing)
- iOS Simulator (Mac only) or Android Emulator

Project Context

Academic Information

- **Course:** Senior Design 2025
- **Domain:** Nutrition Informatics + Machine Learning
- **Team:** 4-person team (Yuvaraj Vagula, Nhat Le, Rached Arda, Matthew Lam)
- **Duration:** Academic year 2024-2025

Learning Objectives Achieved

1.  Production-ready REST API with FastAPI
2.  Vector similarity search with pgvector
3.  ML model integration in backend pipeline
4.  Cross-platform mobile development
5.  Scalable database design with migrations
6.  Software engineering best practices (testing, docs, version control)
7.  Docker containerization and orchestration

Technical Innovations

- **Hybrid Retrieval:** Embedding-based similarity + confidence scoring
- **Semantic Search:** Handles typos, synonyms, and natural language queries
- **Portion Scaling:** Automatic nutrition adjustment based on calories
- **Mixture Aggregation:** Combines multiple similar dishes for better accuracy
- **Multi-tier Architecture:** Clean separation of concerns (mobile → API → database)

 Project Status & Roadmap

 Completed Features

Backend:

-  PostgreSQL schema with pgvector support (7 tables — see Database Schema section)
-  FastAPI with 4 routers (label, dishes, meal-logs/feedback, vision)
-  4-stage retrieval pipeline (retrieval → mixture → scaling → confidence)
-  Full FDA-style nutrition output: 4 required macros + up to 12 nullable micronutrients
-  Automatic meal_log creation on POST /label when device_id provided
-  POST /meal-logs — explicit meal log save endpoint
-  GET /meal-logs/{device_id} — retrieve up to 50 recent entries
-  POST /vision/estimate — camera prediction + auto-write to vision_estimates
-  POST /vision/feedback — user correction submission
-  GET /vision/personalization/{user_id} — personalization profile
-  Database migrations with Alembic
-  Comprehensive data ingestion (2,400,000+ dishes)
-  Unit & integration tests
-  Docker containerization

Mobile:

-  React Native app with Expo

- ☒ 5-tab navigation (Label, Camera, History, Explore, Profile)
- ☒ Label generation UI with full FDA nutrition card + confidence display
- ☒ Barcode scanner screen (BarcodeScannerScreen in Label tab)
- ☒ Camera capture interface for meal photos
- ☒ Daily food tracking (DailyConsumerScreen — client-side FoodContext)
- ☒ Manual entry screen (AddEntryScreen)
- ☒ Platform-aware API configuration
- ☒ TypeScript type safety
- ☒ Material Design UI (React Native Paper)
- ☒ Dark/light theme support (ThemeContext) with 5 accent colour options (amber, emerald, sky, purple, rose)
- ☒ Theme-aware NutritionLabelCard (dark/light FDA label)

Data:

- ☒ 2,400,000+ dishes in database (USDA full + fast food + Open Food Facts)
- ☒ 2,400,000+ searchable variants with embeddings
- ☒ HNSW indexes for fast retrieval
- ☒ Fast food coverage (McDonald's, Burger King, Subway, Taco Bell, etc.)
- ☒ USDA Foundation Foods, SR Legacy, and branded foods coverage

Documentation:

- ☒ Complete setup guide (QUICKSTART.md)
- ☒ Repository structure guide (this file)
- ☒ API documentation
- ☒ Mobile integration guide
- ☒ Dataset expansion plan
- ☒ Final project report (NutriLabelAI_Final_Report.docx)

In Progress

- ☐ History date-range filtering and pagination (GET /meal-logs currently returns newest 50 only)
- ☐ Daily calorie totals endpoint (currently computed client-side in FoodContext)
- ☐ History date-range filtering and pagination (GET /meal-logs currently returns newest 50 only)
- ☐ Explore tab improvements (browse by category, filters)
- ☐ Profile tab enhancements (preferences, goals tracking)
- ☐ Vision model optimization (faster inference, better accuracy)
- ☐ Expanded database (target: 800+ dishes)
- ☐ Additional test coverage for vision pipeline

Future Enhancements

Short-term:

- ☐ User authentication
- ☐ Meal history with daily analytics
- ☐ Favorite dishes
- ☐ Dish categorization in Explore tab

- ☐ Export nutrition data (CSV/PDF)

Medium-term:

- ☐ OCR for nutrition labels (extract from product packaging)
- ☐ Barcode scanning integration
- ☐ Meal planning features
- ☐ Recipe suggestions based on nutrition goals
- ☐ Social sharing and meal comparisons
- ☐ Improved depth sensing accuracy

Long-term:

- ☐ AI-powered meal recommendations
- ☐ Integration with fitness trackers
- ☐ Restaurant menu integration (API partnerships)
- ☐ Multi-language support
- ☐ Web application

Contributing

Development Workflow

1. **Fork the repository**
2. **Create a feature branch:** `git checkout -b feature/your-feature`
3. **Make changes and test:** `pytest app/tests/`
4. **Commit changes:** `git commit -m "feat: add your feature"`
5. **Push to branch:** `git push origin feature/your-feature`
6. **Create pull request**

Code Standards

- **Python:** PEP 8, type hints, docstrings
- **TypeScript:** ESLint + Prettier configuration
- **SQL:** Normalized schema, indexed foreign keys
- **Git:** Conventional commit messages
- **Tests:** Required for new features

Commit Message Convention

```
feat: add new feature
fix: bug fix
docs: documentation changes
style: formatting, missing semicolons, etc.
refactor: code restructuring
test: adding tests
chore: maintenance tasks
```

Project Information

Repository: github.com/yvagula06/Senior-Design-2025

Primary Developer: @yvagula06

Academic Year: 2024-2025

Tech Stack: React Native + FastAPI + PostgreSQL + pgvector

License: MIT (or specify your license)

Documentation Index

Getting Started (Read First)

1. [QUICKSTART.md](#) - Complete setup walkthrough (5-10 minutes)
2. [REPO_BREAKDOWN.md](#) - This file: full repository guide
3. [README.md](#) - Project overview and quick reference

Technical Documentation

- [LABEL_ROUTER_API.md](#) - API endpoint documentation
- [MOBILE_INTEGRATION.md](#) - Mobile-backend integration
- [DATASET_PLAN.md](#) - Data acquisition and expansion strategy
- [schema.sql](#) - PostgreSQL database schema

Academic Reports (Reports/)

- [Reports/NutriLabelAI_Final_Report.docx](#) - Final project report ★ ★
- [Reports/NutriLabelAI_Final_Report_REVISED.md](#) - Revised final report (Markdown)
- [Reports/MIDTERM_REPORT.md](#) - Academic midterm report
- [Reports/Senior Design Fall 2025 Proposal.pdf](#) - Project proposal

Development Guides

- [mobile/README.md](#) - Mobile app architecture
- [scripts/README_PREPROCESSING.md](#) - Data preprocessing

API Reference

- **Swagger UI:** <http://localhost:8000/docs> (interactive API docs)
 - **ReDoc:** <http://localhost:8000/redoc> (alternative documentation)
-

Project Highlights

What Makes This Project Stand Out

1. Semantic Search at Scale

- 384-dimensional embeddings with HNSW indexing
- Sub-100ms query latency for 2.4M+ dishes
- Handles natural language variations, typos, and synonyms

2. Dual-Mode Estimation

- Text-based: Retrieval → Mixture → Scaling → Confidence
- Vision-based: Segmentation → Volume → Classification → Mapping
- Each stage adds intelligence and reliability
- Transparent confidence scoring for both modes

3. Production-Ready Architecture

- Docker containerization for consistent environments
- Database migrations for schema versioning
- Comprehensive test suite
- Type-safe TypeScript + Python with Pydantic

4. Mobile-First Design

- Native iOS & Android with single codebase
- Instant development with Expo Go
- Material Design for polished UI
- Platform-aware API configuration

5. Comprehensive Documentation

- 11 detailed markdown guides
- Interactive API documentation
- Step-by-step setup instructions
- Architecture diagrams and workflows
- Full academic report suite (proposal, midterm, final)

Technical Achievements

-  2,406,474 dishes with 2,406,474 searchable variants
-  pgvector HNSW indexes for fast similarity search
-  95%+ confidence on exact matches
-  Camera-based meal estimation with multiple modes
-  Vision pipeline with segmentation and volume estimation
-  Type-safe end-to-end (TypeScript + Pydantic)
-  Fully containerized deployment
-  Cross-platform mobile app with camera integration
-  RESTful API with OpenAPI documentation
-  Production-grade error handling and validation

Tips for New Contributors

First-Time Setup

1. Read [QUICKSTART.md](#) for environment setup
2. Start backend and populate database
3. Test API: `curl http://localhost:8000/docs`

4. Run mobile app and test both label generation and camera capture
5. Explore codebase starting with [app/main.py](#) and [mobile/App.tsx](#)

Understanding the Codebase

- **Start with:** [app/api/label_router.py](#) (text-based endpoint)
- **Vision pipeline:** [app/api/vision_router.py](#) (camera-based endpoint)
- **Then explore:** [app/services/](#) (business logic for both pipelines)
- **Mobile entry:** [mobile/src/screens/Label/LabelHomeScreen.tsx](#)
- **Camera interface:** [mobile/src/screens/Vision/CameraCaptureScreen.tsx](#)
- **Database:** [app/db/models.py](#) and [schema.sql](#)

Common Tasks

- **Add a dish:** Use [scripts/insert_test_data.py](#)
- **Query database:** Use [scripts/inspect_db.py](#)
- **Test API:** Use FastAPI docs at [/docs](#) endpoint
- **Debug mobile:** Check Network tab in Expo Go developer menu
- **View logs:** `docker logs nutrition_api -f`

Getting Help

- Check [QUICKSTART.md](#) troubleshooting section
- Review API documentation at [/docs](#)
- Inspect database with provided scripts
- Check GitHub issues for similar problems

Last Updated: April 23, 2026

Version: 1.2 - Barcode Scanner, Theme Support, Expanded Scripts & Reports